

a3-Pretrained-Language-Models-dan-jang

November 16, 2023

1 CS410: Natural Language Processing, Fall 2023

1.1 A3: Pretrained Language Models, Dan Jang - 11/12/2023

Description of Assignment

Introduction Using the same training & testing datasets from our first & second assignments, in this assignment, *A3: Pretrained Language Models*, we will be exploring & comparing the performance of two, specific **Pretrained Language Models (PLMs)**:

1. *Google AI's BERT (Bidirectional Encoder Representations from Transformers)*

&

2. *OpenAI's GPT-2 (Generative Pretrained Transformer 2).*

Like the previous two assignment, this assignment focuses on implementing a text-classification model that predicts sentiment & the same training/testing datasets, where in *A3* specifically, we (where, in comparison to our text-classification model approach in our previous *A2* assignment, we used **pretrained embeddings** from the **Word2Vec** and **GloVe models** instead).

Data Preparation Like our previous two assignments, we will use a pair of training & testing datasets containing product customer reviews, which is named the “*Multilingual Amazon Reviews Corpus*”, in a `.json` container format, with several columns. The assignment will focus on a smaller subset of the original dataset, where we will focus on **two (2) columns**: 1. “review_title” - self-explanatory 2. “stars” - an integer, either 1 or 5, where the former indicates “negative” and 5 indicates “positive.”

There will be a training set & a test set.

We will load the dataset using Python & use respective libraries to implement our text-classification model.

In contrary to the last two assignments, there will be no preprocessing done manually, except in using each *PLM's* specific tokenizers to prepare our data for each run of our text-classification model implementations.

Implementation of Pretrained Language Models (PLMs) We will use *HuggingFace* libraries (e.g. `transformers`) to access, then correspondingly experiment with the *BERT* and *GPT-2 pretrained language models (PLMs)*, focusing on these aspects:


1. **Comparison** of the two pretrained language models.

2. Model Evaluation, Results, & Analysis (and comparison) of our two *PLMs*.

BERT (Bidirectional Encoder Representations from Transformers) Pretrained Language Model (PLM) The first *PLM* we will be using is [BERT](#) (*Bidirectional Encoder Representations from Transformers*), which was first described in a [paper](#) published in May 24th, 2019, by the [Google AI Language Team](#), authored by Researchers Jacob Devlin, Ming-Wei Chang, Kenton Lee, and Kristina Toutanova.

[BERT](#) (*Bidirectional Encoder Representations from Transformers*) was created-in-mind, to be a model designed particularly for training “*their own state-of-the-art question answering system*” - built upon previous natural language processing research in pretraining contextual representations, e.g. *Semi-Supervised Sequence Learning*, *Generative Pre-Training*, etc. ([Google AI Language, Devlin & Chang](#), November 2 2018).

Moreover, [BERT](#) (*Bidirectional Encoder Representations from Transformers*), at the time, this model was the “*first **deeply** **directional**, **unsupervised** language representation, pretrained using only a plain text corpus*”, using *Wikipedia* as its training source ([2018](#)).




OpenAI’s GPT-2 (Generative Pretrained Transformer) Pretrained Language Model (PLM) An almost infamously named *PLM* model, [GPT-2](#) (*Generative Pretrained Transformer 2*) is the second predecessor of the **ChatGPT** model that is known popularly today (as of November 10th, 2023, **ChatGPT** utilizes specifically for its *PLM* models, **GPT-3.5 Turbo** for free & **GPT-4+** [/w [Vision](#)] for premium users), our second *PLM* [OpenAI](#)’s [GPT-2](#) (*Generative Pretrained Transformer 2*).

On February 14th, 2019, [OpenAI](#) describes first in a [blog post](#) & [technical paper](#), their announcement & description of their new “*large-scale unsupervised language model*”, named **GPT-2** ([OpenAI](#), February 14 2019).

In **GPT-2**’s [technical paper](#), [OpenAI](#) researchers Alec Radford, Jeffrey Wu, Rewon Child, David Luan, Dario Amodei, and Ilya Sutskever describes the **GPT-2** model - a **unidirectional** model trained on a dataset of *8 million web pages* with **1.5 billion parameters** ([OpenAI](#), [Radford et al.](#), February 14 2019).

Comparison of the BERT vs. GPT-2 Model Architectures & Designs While both [BERT](#) and [GPT-2](#) share great similarities in their overall architecture & design, one of the biggest differences between these two *PLMs* lies in the *deeply bidirectional* nature of [BERT](#) & the *unidirectional* nature of [GPT-2](#). Specifically, BERT had been traditionally used preferably over GPT-2 in certain natural language processing (NLP) tasks, as it seemed to be designed with a more efficient model design & architecture ([Kehlbeck, et al.](#), July 31 2021).

Furthermore, the nature of embedding spaces in *BERT* and *GPT-2* differ greatly, as shown below of the first layers of each respective *PLM*:



Most importantly, with regards to the strengths of our two *PLMs*, *BERT* may be better for tasks surrounding contextual understanding of sentences, while *GPT-2* may be more geared towards tasks related to generating text & general language tasks (Fhal, January 17 2023).

Text Classification Model To build our text-classification model, we will follow these steps:

1. Implementing & setting up our two (2) **Pretrained Language Models (PLMs)**, *BERT* and *[GPT-2](C:)*.
2. Using the specific tokenizer used for each respective *PLM* to prepare the training & testing text data.
3. Training of the text-classification model using the specifically-tokenized training dataset for each *PLM*, based on “sentiment_train.json.”
4. Evaluation of our text-classification model using the specifically-tokenized testing dataset for each *PLM*, based on “sentiment_test.json.”

Results & Analysis A detailed analysis of the model’s performance by comparing the results from the output of our two algorithms, where we will include the following:

1. *F1-score* or other relevant metrics.
2. Any challenges or limitations of the text-classification model/task.

Additionally, we will also try to provide a comparative analysis based on the results from our two previous assignments, our first assignment, *A1: Sentiment Analysis Text Classification* & from our second assignment, *A2: Word2Vec and GloVe Embeddings*.

Specifically, we recall from our previous two assignments, that we first implemented text-classification models based on two suitable algorithms (*A1*), and in our second, implemented the usage of *Word2Vec* & *GloVe* pretrained embeddings to assist in our text-classification tasks (*A1*).

Requirements

1.1.1 Libraries & Constants Initialization

```
[1]: ##### CS410: Natural Language Processing, Fall 2023 - 11/13/2023
##### A3: Pretrained Language Models (PLMs), Dan Jang - Initializations:
↳ Libraries, Models, Data, & Constants

### 0.) Libraries
```

```

from sklearn.naive_bayes import GaussianNB
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score
from sklearn.metrics import confusion_matrix
from sklearn.metrics import f1_score
from sklearn.metrics import classification_report
from sklearn.model_selection import train_test_split
from transformers import AutoTokenizer, AutoModel
import torch
import json
import pandas
#import huggingface_hub
import numpy as np
import random

## 1.0.) Constants, Variables, & Datasets

percentnness = float(100)
# trainfile = str(trainfile)
# testfile = str(testfile)
traindata = []
testdata = []

# Loading the pretrained language models (PLMs), BERT & GPT-2 using HuggingFace
↳libraries!
## Bertie -> BERT PLM
## Bumblebee -> GPT(transformer)-2 PLM
print("Downloading / Loading the BERT Pretrained Language Model (PLM) through
↳HuggingFace libraries...")
bertie = AutoModel.from_pretrained("bert-base-uncased")
print("Loading the BERT Specific Tokenizer...")
bertie_tokens = AutoTokenizer.from_pretrained("bert-base-uncased")
print("...the BERT (Base, pretrained, uncased tokenization) Pretrained Language
↳Model (PLM) has been downloaded / loaded!\n")

print("Downloading / Loading the GPT-2 Pretrained Language Model (PLM) through
↳HuggingFace libraries...")
bumblebee = AutoModel.from_pretrained("gpt2")
print("Loading the GPT-2 Specific Tokenizer...")
bumblebee_tokens = AutoTokenizer.from_pretrained("gpt2")
print("...the GPT-2 Pretrained Language Model (PLM) has been downloaded /
↳loaded!\n")
print("Checking for GPT-2 tokenizer padding token...")
if bumblebee_tokens.pad_token is None:
    print("GPT-2 tokenizer has no padding token!")
    print("...setting GPT-2 tokenizer padding token...")
    bumblebee_tokens.pad_token = bumblebee_tokens.eos_token

```

```
print("...GPT-2 tokenizer padding token set!")
print("\n\nModel initialization all done!\n")
```

Downloading / Loading the BERT Pretrained Language Model (PLM) through HuggingFace libraries...
 Loading the BERT Specific Tokenizer...
 ...the BERT (Base, pretrained, uncased tokenization) Pretrained Language Model (PLM) has been downloaded / loaded!

Downloading / Loading the GPT-2 Pretrained Language Model (PLM) through HuggingFace libraries...
 Loading the GPT-2 Specific Tokenizer...
 ...the GPT-2 Pretrained Language Model (PLM) has been downloaded / loaded!

Checking for GPT-2 tokenizer padding token...
 GPT-2 tokenizer has no padding token!
 ...setting GPT-2 tokenizer padding token...
 ...GPT-2 tokenizer padding token set!

Model initialization all done!

1.1.2 Main Implementation: *Text Classification, with data-processed using respective tokenizers from \mathcal{E} with Two (2) Pretrained Language Models, BERT and GPT-2.*

```
[2]: ##### CS410: Natural Language Processing, Fall 2023 - 11/13/2023
##### A3: Pretrained Language Models (PLMs), Dan Jang - Main Implementation
##### Objective: Exploring Natural Language Processing (NLP), by building a
↳text-classifier
#### for a text classification task, predicting whether a piece of text is
↳"positive" or "negative."
#### ...focusing on two (2) pretrained language models (PLMs),
#### ...BERT (Bidirectional Encoder Representations from Transformers) &
↳OpenAI's GPT-2 (Generative Pretrained Transformer),
#### ...and using the respective tokenizers to each PLM to perform the
↳text-classification task as aforementioned

### 1.1.a) Logistic Regression algorithm using sklearn.linear_model.
↳LogisticRegression
### https://scikit-learn.org/stable/modules/generated/sklearn.linear_model.
↳LogisticRegression.html
### Returns four (4) things:
# I.) accuracy_score,
# II.) f1_score,
# III.) confusion_matrix,
# & IV.) classification_report.
```

```

def plm(model, xtrain, ytrain, xtest, ytest):

    lreg = LogisticRegression()

    lreg.fit(xtrain, ytrain)
    predictionresults = lreg.predict(xtest)

    return accuracy_score(ytest, predictionresults), f1_score(ytest,
    ↪predictionresults), confusion_matrix(ytest, predictionresults),
    ↪classification_report(ytest, predictionresults)

### A3-Specific: Implementation functions for implementing the two pretrained
    ↪language models, BERT & GPT-2.
## Pretrained Language Model (PLM) Tokenizer Implementation Function
def plmodel(words, model, tokenizer):
    if tokenizer.pad_token is None:
        raise ValueError("[Debug A3.1 - PLModel()]: Tokenizer has no padding
    ↪token for the current model!")
    wordlist = tokenizer(words, return_tensors="pt", truncation=True,
    ↪padding=True)

    with torch.no_grad():
        results = model(**wordlist)

    return results.last_hidden_state.mean(dim=1).squeeze().numpy()
    # wordlist = words.split()
    # vecs = [model[w] for w in wordlist if w in model]
    # if vecs:
    #     return sum(vecs) / len(vecs)
    # else:
    #     return [0] * model.vector_size

# ## GPT(transformer)-2 Model Implementation Function
# def bumblebee(words, model):
#     beds = [avgwordvec(w, model) for w in words]
#     return beds

def main(): #trainfile, testfile):
    print("Welcome, this is the main program for A3: Pretrained Language Models.
    ↪")
    print("Written by Dan J. for CS410: Natural Language Processing, Fall 2023.
    ↪")
    print("\nWe will use two (2) pretrained language models (PLM), BERT & GPT-2.
    ↪\n...to create a text-classifier to guess negative or positive
    ↪sentimentality based on various text-reviews of products.")

```

```

# 1.0.I.A) Debug Statements #1a for dataset loading times:
print("\nLoading the training & testing datasets...")
# with open(trainfile, "r") as trainfile:
with open("sentiment_train.json", "r") as trainfile:
    #traindata = json.load(trainfile)
    for row in trainfile:
        traindata.append(json.loads(row))

trainframe = pandas.DataFrame(traindata)

# with open(testfile, "r") as testfile:
with open("sentiment_test.json", "r") as testfile:
    #testdata = json.load(testfile)
    for row in testfile:
        testdata.append(json.loads(row))

testframe = pandas.DataFrame(testdata)

# 1.0.I.B) Debug Statements #1b for dataset loading times:
print("Successfully loaded the training & testing datasets!\n")

## 1.0.1.) Initial Preprocessing of the training & testing data
## First, we isolate our two (2) columns, "review_title" & "stars."
## Second, we will convert values in the "stars" column so that 1
↪ [negative] = 0 & 5 [positive] = 1.
## This will allow us to make the negative or positive sentiment a binary
↪ value-based thingy.
trainframe = trainframe[['review_title', 'stars']]
trainframe['stars'] = trainframe['stars'].apply(lambda x: 1 if x == 5 else
↪ 0)

testframe = testframe[['review_title', 'stars']]
testframe['stars'] = testframe['stars'].apply(lambda x: 1 if x == 5 else 0)

## A3-Specific: From our Slack channel (#nlp_f23), using tip to only use
↪ 25% of training dataset, evenly split
## Credits to Classmate Will McIntosh from the Slack thread started by
↪ classmate Saurav Kumar Singh
## & also, full credits to classmate Will McIntosh for the following code
↪ for GPU usage:

#### Credits to Will McIntosh (11/11/2023):
# Testing
print(f"Is a GPU available: {torch.cuda.is_available()}")
#print(f"Is this instance using a GPU?: {next(model.parameters()).is_cuda}")
#### From Slack, #nlp_f23.

```

```

trainframe = trainframe.sample(frac=1, random_state=69)
trainframe = trainframe.iloc[:int(0.25 * len(trainframe))]

# y2train = trainframe['stars']
# print("[A3 Debug Size-Print #3] y2train", len(y2train))

ytest = testframe['stars']
# print("[A3 Debug Size-Print #4] ytest", len(ytest))

## Evenly split frames
x3train1, x3train2 = train_test_split(trainframe, test_size=0.5,
↪random_state=69)

y3train1 = x3train1['stars']
y3train2 = x3train2['stars']

#print("[A3 Debug Size-Print #1] x3train1 & x3train2", len(x3train1),
↪len(x3train2))
## A3-Specific: Applying BERT & GPT-2 PLM Specific Tokenization
print("Tokenizing the training & testing datasets for BERT...")
x2train1 = x3train1['review_title'].apply(lambda x: plmodel(x, bertie,
↪bertie_tokens))
#x2train1 = trainframe['review_title'].apply(lambda x: plmodel(x, bertie,
↪bertie_tokens))
x2test1 = testframe['review_title'].apply(lambda x: plmodel(x, bertie,
↪bertie_tokens))
print("BERT Tokenization has been applied to its training & testing
↪datasets!")

#print("[A3 Debug Size-Print #2a] x2train1 & x2test1", len(x2train1),
↪len(x2test1))

print("Tokenizing the training & testing datasets for GPT-2...")
x2train2 = x3train2['review_title'].apply(lambda x: plmodel(x, bumblebee,
↪bumblebee_tokens))
#x2train2 = trainframe['review_title'].apply(lambda x: plmodel(x,
↪bumblebee, bumblebee_tokens))
x2test2 = testframe['review_title'].apply(lambda x: plmodel(x, bumblebee,
↪bumblebee_tokens))
print("GPT-2's tokenization has been applied to its training & testing
↪datasets!")

#print("[A3 Debug Size-Print #2b] x2train2 & x2test2", len(x2train2),
↪len(x2test2))

```



```

### 1.0.2b) Run Text-Classification Algorithms & Print the Model Results
print("-----\n")
print("Running text-classification model le training & testing datasets_
↳(with pretrained language models, BERT & GPT-2)...")

print("Running Logistic Regression algorithm, version A.) BERT...")
bed1accuracy, bed1f1, bed1cmatrix, bed1creport = plm(bertie, x2train1.
↳tolist(), y3train1, x2test1.tolist(), ytest)
print("..The data processing from our first PLM, BERT, is done!")

print("Running Logistic Regression algorithm, version B.) GPT-2...")
bed2accuracy, bed2f1, bed2cmatrix, bed2creport = plm(bumblebee, x2train2.
↳tolist(), y3train2, x2test2.tolist(), ytest)
print("..The data processing from our second PLM, GPT-2, is done!")

print("...All Done!")
print("-----\n")

print("Here are le results [Logistic Regression, with comparative results_
↳between two (2) PLMs, BERT & GPT-2]...\n")
print("Logistic Regression Algorithm, Version A: BERT Pretrained Language_
↳Model-based Text-Classification Performance, Metrics, & Results:")
print("...Accuracy was found to be, ", bed1accuracy * percentness, "%,")
print("...F1 Score was found to be: ", bed1f1, ",")
print("...with a Confusion Matrix: \n", bed1cmatrix, ",")
print("...& lastly, the classification Report: \n", bed1creport)
print("-----\n")

print("Logistic Regression Algorithm, Version B: GPT-2 Pretrained Language_
↳Model-based Text-Classification Performance, Metrics, & Results:")
print("...Accuracy was found to be, ", bed2accuracy * percentness, "%,")
print("...F1 Score was found to be: ", bed2f1, ",")
print("...with a Confusion Matrix: \n", bed2cmatrix, ",")
print("...& lastly, the classification Report: \n", bed2creport)
print("-----\n")

if __name__ == "__main__":
    main()

```

Welcome, this is the main program for A3: Pretrained Language Models.
 Written by Dan J. for CS410: Natural Language Processing, Fall 2023.

We will use two (2) pretrained language models (PLM), BERT & GPT-2.
 ...to create a text-classifier to guess negative or positive sentimentality
 based on various text-reviews of products.

```

Loading the training & testing datasets...
Successfully loaded the training & testing datasets!

Is a GPU available: False
Tokenizing the training & testing datasets for BERT...
BERT Tokenization has been applied to its training & testing datasets!
Tokenizing the training & testing datasets for GPT-2...
GPT-2's tokenization has been applied to its training & testing datasets!
-----

Running text-classification model le training & testing datasets (with
pretrained language models, BERT & GPT-2)...
Running Logistic Regression algorithm, version A.) BERT...

c:\tools\miniconda3\lib\site-packages\sklearn\linear_model\_logistic.py:460:
ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:
    https://scikit-learn.org/stable/modules/preprocessing.html
Please also refer to the documentation for alternative solver options:
    https://scikit-learn.org/stable/modules/linear_model.html#logistic-
regression
    n_iter_i = _check_optimize_result(

..The data processing from our first PLM, BERT, is done!
Running Logistic Regression algorithm, version B.) GPT-2...
..The data processing from our second PLM, GPT-2, is done!
...All Done!
-----

Here are le results [Logistic Regression, with comparative results between two
(2) PLMs, BERT & GPT-2]...

Logistic Regression Algorithm, Version A: BERT Pretrained Language Model-based
Text-Classification Performance, Metrics, & Results:
...Accuracy was found to be,  91.55 %,
...F1 Score was found to be:  0.9158785465405676 ,
...with a Confusion Matrix:
[[911  89]
 [ 80 920]] ,
...& lastly, the classification Report:

```

	precision	recall	f1-score	support
0	0.92	0.91	0.92	1000
1	0.91	0.92	0.92	1000
accuracy			0.92	2000
macro avg	0.92	0.92	0.92	2000

```
weighted avg      0.92      0.92      0.92      2000
```

Logistic Regression Algorithm, Version B: GPT-2 Pretrained Language Model-based Text-Classification Performance, Metrics, & Results:

...Accuracy was found to be, 87.4 %,

...F1 Score was found to be: 0.872598584428716 ,

...with a Confusion Matrix:

```
[[885 115]
```

```
[137 863]] ,
```

...& lastly, the classification Report:

	precision	recall	f1-score	support
0	0.87	0.89	0.88	1000
1	0.88	0.86	0.87	1000
accuracy			0.87	2000
macro avg	0.87	0.87	0.87	2000
weighted avg	0.87	0.87	0.87	2000

c:\tools\miniconda3\lib\site-packages\sklearn\linear_model_logistic.py:460:

ConvergenceWarning: lbfgs failed to converge (status=1):

STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

1.1.3 A3: Pretrained Language Models (PLMs), Text-Classification Model Performance Analysis & Discussion

Initial Data Results, Metrics, & Analysis Using the Logistic Regression Algorithm like previously for A2, this algorithm was used to implement our text-classification model, for the text-classification task in both *Version A* & *B*.

As shown below, *Version A* featured *Google AI's BERT* model & *Version B* featured *OpenAI's GPT-2* model - where we saw the following results:

Version A - BERT Pretrained Language Model (PLM) Results:

Accuracy: ~91.6%

F1 Score: 0.9158785465405676

Confusion Matrix:

```
[[911  89]
 [ 80 920]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.92	0.91	0.92	1000
1	0.91	0.92	0.92	1000
accuracy			0.92	2000
macro avg	0.92	0.92	0.92	2000
weighted avg	0.92	0.92	0.92	2000

Version B - GPT-2 Pretrained Language Model (PLM) Results:

Accuracy: 87.4%

F1 Score: 0.872598584428716

Confusion Matrix:

```
[[885 115]
 [137 863]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.87	0.89	0.88	1000
1	0.88	0.86	0.87	1000
accuracy			0.87	2000
macro avg	0.87	0.87	0.87	2000
weighted avg	0.87	0.87	0.87	2000

Pretrained Language Model (PLM) Comparative Analysis & Discussion From the results above, we see that both *Google AI's BERT* & *OpenAI's GPT-2 PLMs* fared very well in accuracy, of ~91.6% & 87.4% respectively.

Between these two *PLMs*, *BERT* & *GPT-2*, we can see that from our performance results that there is/was a slight, low *difference of accuracy*, where *BERT* has a *higher accuracy* of ~4.2% in *comparison to GPT-2*.

Also, by looking at the respective Confusion Matrices for our two *PLMs*, we can see that *BERT*, of course, was able to classify sentiment more accurately than *GPT-2*, but we also can observe that *BERT* seems to be *a little bit* more efficient in detecting *negative* sentimentality vs. *GPT-2* lean towards *positive* sentimentality - however, *GPT-2's* decreased accuracy in detecting *negative* sentiment seems to be more significant in magnitude compared to the aforementioned, *negative* over *positive* sentimentality bias, of *BERT*.

Interestingly, a possible scenario where this supposed sentimentality bias in our text-classification task may present itself as to yield a different result, at least for accuracy & the F1 score, may lie in how the positive & negative sentiment composition during the processing of our training & testing dataset, particularly, at the time of randomized splitting. Specifically, if we had somehow chosen a random seed in our splitting code, that supposedly had much more positive or negative sentimentality in the composition of the reviews, we could have seen variations from our initial performance & results as described above.

Ergo, between *Google AI*'s *BERT* & *OpenAI*'s *GPT-2 PLMs*, *BERT* seems to be a clearly more suitable model, at least for the text-classification task & model with our assignment's training & dataset, "*Multilingual Amazon Reviews Corpus*."

Analysis of Current & Previous Text-Classification Models & Performance Results from A1 & A2 In comparison to the performance seen from A2: *Word2Vec & GloVe Embeddings*, both *Google AI*'s *BERT* & *OpenAI*'s *GPT-2 PLMs* outperformed the accuracies & F1 scores from either Word2Vec or GloVe pretrained embeddings - where we recall the A2 results as follows:

A2: Version A, *Word2Vec* Model:

Accuracy: ~86.3%,
F1 Score: 0.86105476673428

Confusion Matrix:

```
[[1754 246]
 [ 302 1698]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.85	0.88	0.86	2000
1	0.87	0.85	0.86	2000
accuracy			0.86	4000
macro avg	0.86	0.86	0.86	4000
weighted avg	0.86	0.86	0.86	4000

A2: Version B, *GloVe* Model:

Accuracy: ~69.69%,
F1 Score: 0.7313829787234042

Confusion Matrix:

```
[[1138 862]
 [ 350 1650]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.76	0.57	0.65	2000
1	0.66	0.82	0.73	2000

accuracy			0.70	4000
macro avg	0.71	0.70	0.69	4000
weighted avg	0.71	0.70	0.69	4000

Although, we do see only a *slight* ~1.1% **accuracy** improvement when we compare results from [Word2Vec](#) vs. [GPT-2](#).

However, this is a neat & expected observation, as we can recall that both [Word2Vec](#) & [GPT-2](#) are suitable or otherwise geared towards the generation of text, rather than the specificity of text-classification required for efficiency & performance with our Amazon review sentimentality prediction task. As expected, contrarywise, as the results from the implementation of [GloVe](#) pre-trained embeddings & the [BERT PLM](#) show higher performance & accuracy compared to the [Word2Vec](#) pretrained embeddings & [GPT-2 PLM](#), we could also make the observation of this supposed higher efficiency & efficacy by recalling that the [GloVe](#) pretrained embeddings were also designed & trained to create an unsupervised learning algorithm, with training that “*performed on aggregated global word-word co-occurrence statistics*” ([Pennington, et al., 2014](#)), which, like the nature of [BERT](#)’s greater suitability towards text-classification tasks, could explain these higher accuracies & better performance as observed from our previous *A2* & most current, *A3* results.

Recalling further, note that from our first assignment, *A1: Sentiment Analysis Text Classification*, we saw the following results from using *two (2) different algorithms* to implement the text-classification model, specifically, the *Logistic Regression* & the *Gaussian Naïve Bayes* Algorithms:

From A1 Results: Version A: Gaussian Naïve Bayes Algorithm:

Accuracy: ~59.2%

F1 Score: 0.3664596273291925

CConfusion Matrix:

[[948 52]

[764 236]]

Classification Report:

	precision	recall	f1-score	support
0	0.55	0.95	0.70	1000
1	0.82	0.24	0.37	1000
accuracy			0.59	2000
macro avg	0.69	0.59	0.53	2000
weighted avg	0.69	0.59	0.53	2000

From A1 Results: Version B: Logistic Regression Algorithm:

Accuracy: 92.7%

F1 Score: 0.9272908366533865

Confusion Matrix:

[[923 77]

[69 931]]

Classification Report:

	precision	recall	f1-score	support
0	0.93	0.92	0.93	1000
1	0.92	0.93	0.93	1000
accuracy			0.93	2000
macro avg	0.93	0.93	0.93	2000
weighted avg	0.93	0.93	0.93	2000

This is interesting, as while it appears that using *PLMs* over *pretrained embeddings* yielded much significant improvements in both accuracy & performance, the results from using only the Logistic Regression Algorithm with no classifiers had yielded a *92.7% accuracy*, which still stands as one of the most accurate results from our three (3) assignments.

However, this observation & other possible variations that can be observed, in performance & results, may be attributed to both the nature of text-tokenization, usage of classifiers, & particularly, that in the implementation of the first text-classification model of our *A1* assignment, I had trained the text-classification model cases with the *full, eighty-thousand (80,000) rows of training data*, while in our current *A3* implementation, the training data for both the *BERT* & *GPT-2 PLMs* were truncated randomly (with a shared random seed of 69) to a quarter-percent (25%) of the original number of rows, or twenty-thousand (20,000) rows of training data.

Text-Classification Challenges & Limitations The initial & biggest challenge to implement these *pretrained language models (PLMs)* was the very significantly resource & time-intensive requirements for the computation, tokenization, & training of the two (2) *PLMs* used.

With the current version of code & implementation, in average, it appeared that *BERT* took around *two-to-eight (2-to-8) minutes*, while *GPT-2* took longer, around *five-to-fifteen (5-to-15) minutes* long. Both, on average, usually took around *fifteen-to-thirty (15-to-30) minutes* to fully complete.

However, in the earlier stages of coding & implementation, particularly before truncating the training & testing datasets to the currently-set percentage of twenty-five percent (25%) of the original amount of rows/lines of text, the average varied wildly, with completion durations taking *hours, upon hours* to complete or were stopped prior to completion due to time constraints.

Furthermore, even post-truncation adjustments, the pretraining of the two *PLMs*, in terms of computational power & resources, were very intensive, which had caused *multiple, full computer crashes*, leading to a lot of processing delays & mental energy to continue implementing.

Discussion for Future Performance & Efficacy Improvements With the issue of computational resources & the prevention of computer crashes, would be, of course, to finally sign-up for the student discounted *Google Cloud* subscription for using high-resource, cloud computing in *Google Colaboratory*, where NLP compatible, GPU-acceleration is available to expedite *PLM*, model training, or otherwise for running intensive NLP code & tasks. As such, this specific idea for future performance & efficacy improvement will be implemented immediately following this current assignment, *A3*, where I will go ahead & attempt to set-up *Google Colaboratory* for use to hopefully, avoid the aforementioned resource & crash pitfalls for *A4* and our *Group Project*, heh.

Furthermore, I should attempt to start early & try to implement for assignments in smaller code-blocks/pieces.

1.1.4 References & Resources

Libraries & Dependencies

numpy
pandas
torch
random

[HuggingFace_hub](#)

[Google AI's BERT](#)

[OpenAI's GPT-2](#)

[HuggingFace's transformers](#)

```
from transformers import AutoModel(s)
from transformers import AutoTokenizer
from transformers import AutoTokenizer.from_pretrained, AutoModel.from_pretrained
sklearn.linear_model.LogisticRegression
sklearn.model_selection.train_test_split
sklearn.metrics.f1_score
sklearn.metrics.accuracy_score
sklearn.metrics.confusion_matrix
sklearn.metrics.classification_report
nbconvert
```

References & Credits *The Multilingual Amazon Reviews Corpus* by Phillip Keung, Yichao Lu, György Szarvas, & Noah A. Smith (October 6th, 2020)

BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding by Jacob Devlin, Ming-Wei Chang, Kenton Lee, & Kristina Toutanova (May 24th, 2019)

Language Models are Unsupervised Multitask Learners by Alec Radford, Jeffrey Wu, Rewon Child David Luan, Dario Amodei, & Ilya Sutskever (*OpenAI*; February 14th, 2019)

Comparing BERT, GPT-2, and GPT-3: A Look at the Pros and Cons of Popular Pre-Trained Language Models by Em Fhal (January 17th, 2023)

GloVe: Global Vectors for Word Representation by Jeffrey Pennington, Richard Socher, & Christopher D. Manning (2014)

Efficient Estimation of Word Representations in Vector Space [Word2Vec Pre-trained Embeddings] by Tomas Mikolov, Kai Chen, Greg Corrado, & Jeffrey Dean (January 16th, 2013)

Credits to GitHub Copilot & ChatGPT for code implementation assistance.

Special Thanks Thanks to fellow classmate *Will McIntosh* for their helpful tips & tricks for resource management particularly for *PLMs* in the current *A3* assignment, from the `#nlp_f23` class-channel on *pdx-cs* Slack, November 11th, 2023.

Extra Stuff

1.1.5 A3 Results: Raw Output from *BERT* & *OpenAI's GPT-2* Pretrained Language Models (PLMs), using the Logistic Regression Algorithm

Here are the results [Logistic Regression, with comparative results between two (2) PLMs, BERT & GPT-2]

Logistic Regression Algorithm, Version A: BERT Pretrained Language Model-based Text-Classification

...Accuracy was found to be, 91.55 %,

...F1 Score was found to be: 0.9158785465405676 ,

...with a Confusion Matrix:

```
[[911 89]
```

```
[ 80 920]] ,
```

...& lastly, the classification Report:

	precision	recall	f1-score	support
0	0.92	0.91	0.92	1000
1	0.91	0.92	0.92	1000
accuracy			0.92	2000
macro avg	0.92	0.92	0.92	2000
weighted avg	0.92	0.92	0.92	2000

Logistic Regression Algorithm, Version B: GPT-2 Pretrained Language Model-based Text-Classification

...Accuracy was found to be, 87.4 %,

...F1 Score was found to be: 0.872598584428716 ,

...with a Confusion Matrix:

```
[[885 115]
```

```
[137 863]] ,
```

...& lastly, the classification Report:

	precision	recall	f1-score	support
0	0.87	0.89	0.88	1000
1	0.88	0.86	0.87	1000
accuracy			0.87	2000
macro avg	0.87	0.87	0.87	2000
weighted avg	0.87	0.87	0.87	2000

1.1.6 A2 Results: Raw Output from Word2Vec & GloVe Embedding Results, using the Logistic Regression Algorithm

Logistic Regression Algorithm, Version A: Word2Vec Pretrained Model-based Embeddings Performance

...Accuracy was found to be, 86.3 %,

...F1 Score was found to be: 0.86105476673428 ,

...with a Confusion Matrix:

[[1754 246]

[302 1698]] ,

...& lastly, the classification Report:

	precision	recall	f1-score	support
0	0.85	0.88	0.86	2000
1	0.87	0.85	0.86	2000
accuracy			0.86	4000
macro avg	0.86	0.86	0.86	4000
weighted avg	0.86	0.86	0.86	4000

Logistic Regression Algorithm, Version B: GloVe Pretrained Model-based Embeddings Performance,

...Accuracy was found to be, 69.69999999999999 %,

...F1 Score was found to be: 0.7313829787234042 ,

...with a Confusion Matrix:

[[1138 862]

[350 1650]] ,

...& lastly, the classification Report:

	precision	recall	f1-score	support
0	0.76	0.57	0.65	2000
1	0.66	0.82	0.73	2000
accuracy			0.70	4000
macro avg	0.71	0.70	0.69	4000
weighted avg	0.71	0.70	0.69	4000

From A1 Results for Reference: Initial Full 80k-Row Processing Results Raw Output

Algorithm #1, Version A: Gaussian N  ive Bayes Performance, Metrics, & Results:

...Accuracy was found to be, 59.199999999999996 %,

...F1 Score was found to be: 0.3664596273291925 ,

...with a Confusion Matrix:

[[948 52]

[764 236]] ,

...& lastly, the classification Report:

	precision	recall	f1-score	support
0	0.55	0.95	0.70	1000
1	0.82	0.24	0.37	1000
accuracy			0.59	2000
macro avg	0.69	0.59	0.53	2000
weighted avg	0.69	0.59	0.53	2000

Algorithm #2, Version A: Logistic Regression Performance, Metrics, & Results:

...Accuracy was found to be, 92.7 %,
 ...F1 Score was found to be: 0.9272908366533865 ,
 ...with a Confusion Matrix:

[[923 77]

[69 931]] ,

...& lastly, the classification Report:

	precision	recall	f1-score	support
0	0.93	0.92	0.93	1000
1	0.92	0.93	0.93	1000
accuracy			0.93	2000
macro avg	0.93	0.93	0.93	2000
weighted avg	0.93	0.93	0.93	2000

Algorithm #1, Version B: Gaussian N  ive Bayes Performance, Metrics, & Results:

...Accuracy was found to be, 59.3 %,
 ...F1 Score was found to be: 0.36899224806201547 ,
 ...with a Confusion Matrix:

[[948 52]

[762 238]] ,

...& lastly, the classification Report:

	precision	recall	f1-score	support
0	0.55	0.95	0.70	1000
1	0.82	0.24	0.37	1000
accuracy			0.59	2000
macro avg	0.69	0.59	0.53	2000
weighted avg	0.69	0.59	0.53	2000

Algorithm #2, Version B: Logistic Regression Performance, Metrics, & Results:

...Accuracy was found to be, 92.80000000000001 %,
...F1 Score was found to be: 0.9281437125748503 ,
...with a Confusion Matrix:

```
[[926 74]
 [ 70 930]] ,
```

...& lastly, the classification Report:

	precision	recall	f1-score	support
0	0.93	0.93	0.93	1000
1	0.93	0.93	0.93	1000
accuracy			0.93	2000
macro avg	0.93	0.93	0.93	2000
weighted avg	0.93	0.93	0.93	2000

[17]: ##### Jupyter Notebook -> PDF Conversion thingy

```
#!pip install nbconvert
```

```
!jupyter nbconvert a3-Pretrained-Language-Models-dan-jang.ipynb --to pdf
```

[NbConvertApp] Converting notebook a3-Pretrained-Language-Models-dan-jang.ipynb to pdf

[NbConvertApp] Writing 95660 bytes to notebook.tex

[NbConvertApp] Building PDF

[NbConvertApp] Running xelatex 3 times: ['xelatex', 'notebook.tex', '-quiet']

[NbConvertApp] Running bibtex 1 time: ['bibtex', 'notebook']

[NbConvertApp] WARNING | b had problems, most likely because there were no citations

[NbConvertApp] PDF successfully created

[NbConvertApp] Writing 3735153 bytes to a3-Pretrained-Language-Models-dan-jang.pdf